

STAT S4240 002, Homework 2

Brian Weinstein (bmw2148)

July 23, 2015

Problem 1: PCA

(a) Column means

```
> apply(rawData, 2, mean)
      x1      x2      x3      x4      x5
6.049104 -8.277221  4.665532  7.914270 62.138753
```

Row means

```
> apply(rawData, 1, mean)
[1] -0.1277116  20.8162864 -8.8984358  25.5999204 -9.7472153
[6] 64.0626702  22.0392371  23.3914888  31.7598224 -13.8680290
...
[91]  1.2105932  21.2145724 -8.4896595  19.0639963  20.9767512
[96]  3.5962333  22.3461063  0.7145014   6.3080005  64.8829556
```

The nonzero column means indicate that each variable isn't centered. In this context the row means are just the average of the coordinates for each observation.

(b) Empirical covariance matrix

	x1	x2	x3	x4	x5
x1	72.96417	-83.90858	53.23708	120.1162	568.4105
x2	-83.90858	110.89101	-63.89570	-115.9430	-817.3388
x3	53.23708	-63.89570	39.60282	83.7386	445.2511
x4	120.11620	-115.94304	83.73860	232.1333	683.5587
x5	568.41046	-817.33884	445.25112	683.5587	6288.8569

The diagonal values tell us the variance of the variable indicated in the column (or equivalently, the row). The off-diagonal elements indicate the covariance between the two variables that intersect at that element.

(c) The eigenvalues and eigenvectors of the empirical covariance matrix `sig`:

```
> eigen(sig)
$values
[1] 6.557348e+03 1.868951e+02 2.038354e-01 9.775594e-04 9.373658e-05

$vectors
      [,1]      [,2]      [,3]      [,4]      [,5]
[1,] 0.09009603 -0.3247102 -0.383470773 0.82286709 0.24957150
[2,] -0.12797842 0.1364755 0.227047683 -0.11412319 0.94890526
[3,] 0.07028767 -0.1941349 0.894987159 0.37278501 -0.13191135
[4,] 0.11077853 -0.9008231 -0.019718518 -0.40719485 0.10024632
[5,] 0.97892389 0.1636064 0.002946326 -0.07133967 0.09921159
```


where the eigenvectors are represented as columns in the `$vectors` component.

Since it's a symmetric matrix, `sig` has the same left eigenvectors as right eigenvectors.

(d) The loadings are the eigenvectors (see part c). The scores are:

```
> data%*%t(evecs)
      [,1]      [,2]      [,3]      [,4]      [,5]
[1,] -25.9233299 -50.96254603 -4.06557021 -8.91350845 -10.7755359
[2,] 13.3064897 13.56908728 6.16049505 0.82185440 4.8621981
[3,] -37.6872799 -93.30323983 -1.02562352 -18.15040050 -16.3848300
...
[98,] -27.0931525 -38.88284377 -8.26671502 -5.26069573 -10.6527232
[99,] -13.1627026 -31.46409161 -1.20277265 -5.83681744 -5.6197025
[100,] 85.7232563 184.73133084 10.16179165 33.54024954 36.1560703
```

(e) Proportion of variance explained



We only need one principal component. PC #1 accounts for 97% of the variance on its own, and including any additional PCs introduces more complexity than it's worth.

- (f) The scores for the new observations:

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	-16.855658	-66.790142	6.776900	-15.3066224	-8.5649268
[2,]	4.861635	26.789815	-3.710359	6.3835809	2.7072847
[3,]	2.536648	1.355434	1.590414	-0.2530266	0.8012362
[4,]	-30.863147	-57.879087	-5.570018	-9.7920953	-12.6610309
[5,]	-11.787862	-14.507653	-3.653059	-1.8812777	-4.6755677

- (g) Coordinates of the projections in the original space:

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	-10.871807	-38.241645	4.665532	7.91427	62.13875
[2,]	37.567018	66.568511	4.665532	7.91427	62.13875
[3,]	16.460691	25.729053	4.665532	7.91427	62.13875
[4,]	5.978194	-41.185647	4.665532	7.91427	62.13875
[5,]	1.857621	-6.945432	4.665532	7.91427	62.13875

Euclidean distance from the original data points.

```
[1] 81.49181
[1] 88.13304
[1] 36.01572
[1] 79.53924
[1] 19.04255
```


- (h) The error vectors are more or less orthogonal to the direction of the first principal component. This is because the error vectors are defined as the direction from the original points to their *orthogonal projections* onto the reduced-dimension space, which is primarily captured by the first PC.

Problem 2: PCA with Yale Faces B

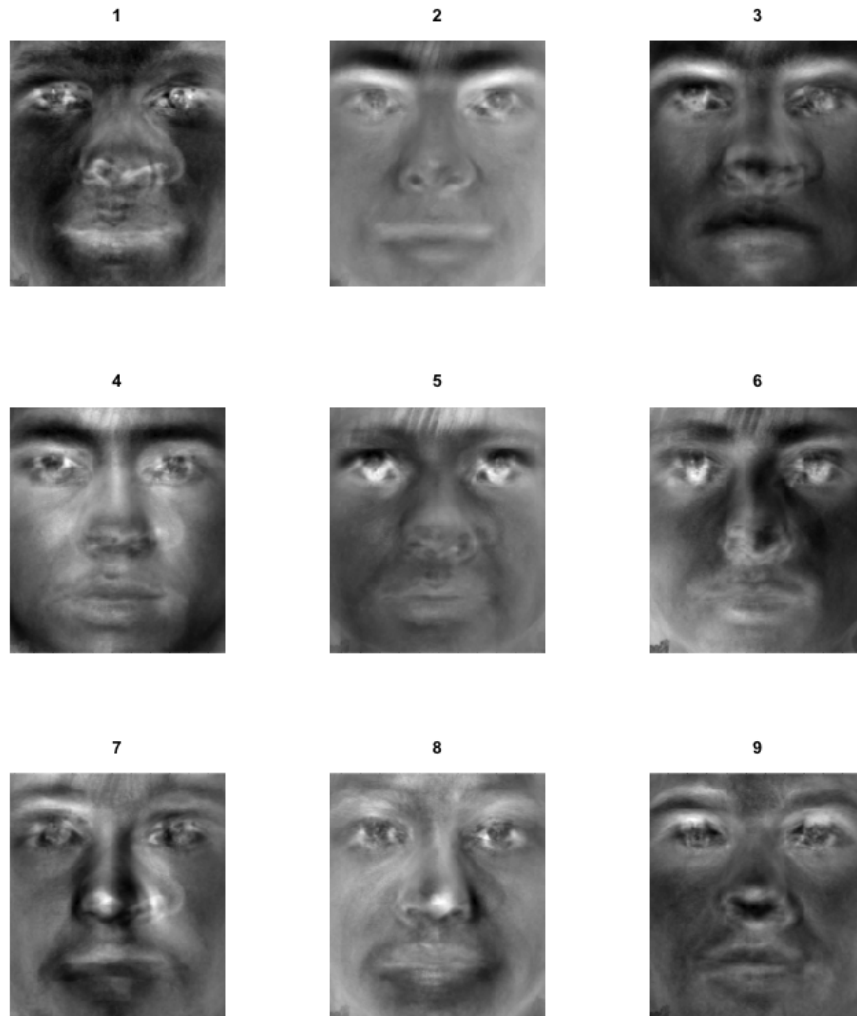
- (a) The matrix is 152 rows by 32,256 columns. Each row of the matrix is one photograph (38 subjects with 4 views each), and each column is a pixel in each image (originally 192 x 168).
- (b) Mean face:



(c) Proportion of variance explained



- (d) The first 9 eigenfaces are the (constructed) images that capture the most variable aspects of the faces in the dataset. The first image amplifies the most variable pixels in the original images, the second image amplifies the most variable pixels in a direction orthogonal to the first principal component, the third the most variable pixels in a direction orthogonal to the first two principal components, etc.



(e) Reconstructing `yaleB05_P00A+010E+00.pgm` using 24 eigenfaces:



Reconstructing yaleB05_P00A+010E+00.pgm using 120 eigenfaces, in steps of 5:



After incorporating 14 or 15 eigenfaces, the person is relatively recognizable

- (f) After removing the images of Subject 1, re-running PCA on the reduced dataset, and then reconstructing a photo of Subject 1 using the principal components; the reconstructed photo (right) has similar features to the original (left), but it's definitely not a recognizable image.

It doesn't look much like the original image because we excluded all photos of Subject 1 before computing the principal components. As such, the features present in the subject's original photos weren't incorporated into either the "mean face" or into the principal components.



Problem 3: James 3.7.3

- (a) iii. For a fixed value of IQ and GPA, males earn more on average than females provided that the GPA is high enough.

$$\text{Salary} = 50 + (20 * GPA) + (0.07 * IQ) + (35 * I_{female}) + (0.01 * GPA * IQ) + (-10 * GPA * I_{female})$$

Thus, once GPA is high enough, the $(-10 * GPA * I_{female})$ term dominates and reduces female salary.

- (b) For a female ($I_{female} = 1$) with $IQ = 110$ and $GPA = 4.0$, the salary (in thousands of dollars) is predicted to be:

$$\text{Salary} = 50 + (20 * 4.0) + (0.07 * 110) + (35 * 1) + (0.01 * 4.0 * 110) + (-10 * 4.0 * 1) = 137.1$$

- (c) False. A small coefficient for the GPA/IQ interaction term does not mean there's little evidence of an interaction effect. To determine how strong the evidence of interaction is, we should look at the p-value of the term's coefficient β_4 . If the term is significant, the small coefficient indicates that the interaction between GPA and IQ has little effect on Salary.

Problem 4: James 3.7.4

- (a) If the true relationship is linear (with a random error term), the training RSS for the cubic regression will be lower than the training RSS for the linear regression. This is because, while using training data, a more flexible model (e.g., cubic vs linear) will always fit better than a less flexible one.
- (b) The testing RSS, on the other hand, will be lower for the linear regression than the testing RSS for the cubic regression. By using the cubic regression on a truly linear relationship, the RSS will indicate that we overfit the data.
- (c) If we know that the relationship isn't linear, but we don't know how far it is from linear, we'll still have a lower RSS for the cubic regression than for the linear regression. As noted in part (a), this is because a more flexible model will always reduce the training RSS more than a less flexible one.
- (d) If we again know that the relationship isn't linear, but we don't know how far it is from linear, it's not clear which model will have a lower test RSS. We need to know if the model is closer to linear or closer to cubic to determine which model would have a lower test RSS.

Problem 5: question

Problem 6: question

Todo list