

STAT S4240 002, Homework 3

Brian Weinstein (bmw2148)

July 30, 2015

Problem 1: Naive Bayes Text Classification: Data Preparation

See hw03_q1.R for code.

(a) Text pre-processing

```
# load functions from hw03.R
source("hw03.R")

# preprocess text
preprocess.directory("datasets/FederalistPapers/fp_hamilton_train")
preprocess.directory("datasets/FederalistPapers/fp_hamilton_test")
preprocess.directory("datasets/FederalistPapers/fp_madison_train")
preprocess.directory("datasets/FederalistPapers/fp_madison_test")
```

(b) Loading the cleaned data

```
hamilton.train <- read.directory("datasets/FederalistPapers/fp_hamilton_train_clean")
hamilton.test <- read.directory("datasets/FederalistPapers/fp_hamilton_test_clean")
madison.train <- read.directory("datasets/FederalistPapers/fp_madison_train_clean")
madison.test <- read.directory("datasets/FederalistPapers/fp_madison_test_clean")
```

(c) Create a dictionary from all of the documents

```
dictionary <- make.sorted.dictionary.df(c(hamilton.train, hamilton.test,
                                         madison.train, madison.test))
```

(d) Creating document-term-matrices for each of the datasets

```
dtm.hamilton.train <- make.document.term.matrix(infiles=hamilton.train,
                                                dictionary=dictionary)
dtm.hamilton.test <- make.document.term.matrix(infiles=hamilton.test,
                                                dictionary=dictionary)
dtm.madison.train <- make.document.term.matrix(infiles=madison.train,
                                                dictionary=dictionary)
dtm.madison.test <- make.document.term.matrix(infiles=madison.test,
                                                dictionary=dictionary)
```

(e) Compute the log probabilities for the dictionary in each of the document datasets

```
mu=1/nrow(dictionary)

logp.hamilton.train <- make.log.pvec(dtm.hamilton.train, mu)
logp.hamilton.test <- make.log.pvec(dtm.hamilton.test, mu)
logp.madison.train <- make.log.pvec(dtm.madison.train, mu)
logp.madison.test <- make.log.pvec(dtm.madison.test, mu)
```

Problem 2: Naive Bayes Function

We first estimate the log priors based on the log of the proportion of training documents attributed to each author.

$$p(\text{author} = \text{author}) = \log \left(\frac{\# \text{ of training documents attributed to } \text{author}}{\text{total } \# \text{ of training documents}} \right)$$

Then, using (1) the log probabilities for the dictionary in a Hamilton-authored document and (2) the log probabilities for the dictionary in a Madison-authored document (as computed in **Problem 1**), we can input a new document-term-matrix and classify each document as belonging to one of the authors.

```
naive.bayes <- function(logp.hamilton.train, logp.madison.train,
                        log.prior.hamilton, log.prior.madison, dtm.test){
  # Performs naive bayes classification
  # Inputs: logp.hamilton.train : vector of log probabilities of words
  #           occurring in the hamilton training data
  #           logp.madison.train : vector of log probabilities of words
  #           occurring in the madison training data
  #           log.prior.hamilton : the log prior of hamilton documents
  #           log.prior.madison : the log prior of madison documents
  #           dtm.test          : a document-term-matrix to classify
  # Output: Classification labels for each document in dtm.test

  # calculate the log posterior probabilities
  log.post.hamilton <- log.prior.hamilton + (dtm.test %*% logp.hamilton.train)
  log.post.madison <- log.prior.madison + (dtm.test %*% logp.madison.train)

  # compare the log posterior probabilities and assign to the author
  # with highest probability
  prediction <- data.frame(logPostHam=log.post.hamilton,
                           logPostMad=log.post.madison)
  prediction$pred <- (log.post.hamilton >= log.post.madison)
  prediction$pred <- gsub(TRUE, "Hamilton", prediction$pred)
  prediction$pred <- gsub(FALSE, "Madison", prediction$pred)

  # return a vector of the predictions
  return(prediction$pred)
}
```

Problem 3: Assessing Model Performance

Using the confusionMatrix function from the caret library:

- **Accuracy:** 63% accurate (% of the test papers that are classified correctly)
- **True Positive Rate:** 100% (Hamilton classified as Hamilton divided by the total amount of testing Hamilton papers)

- **True Negative Rate:** 9% (Madison classified as Madison divided by the total amount of testing Madison papers)
- **False Positive Rate:** 91% (Madison classified as Hamilton divided by the total amount of testing Madison)
- **False Negative Rate:** 0% (Hamilton classified as Madison divided by the total amount of testing Hamilton)

```
> confusionMatrix(data=predictions$pred,
+                 reference=predictions$trueValue,
+                 dnn=c("Prediction", "True Value"),
+                 positive="Hamilton")
```

Confusion Matrix and Statistics

	True Value	
Prediction	Hamilton	Madison
Hamilton	16	10
Madison	0	1

Accuracy : 0.6296
 95% CI : (0.4237, 0.806)
 No Information Rate : 0.5926
 P-Value [Acc > NIR] : 0.427258

Kappa : 0.106
 McNemar's Test P-Value : 0.004427

Sensitivity : 1.00000
 Specificity : 0.09091
 Pos Pred Value : 0.61538
 Neg Pred Value : 1.00000
 Prevalence : 0.59259
 Detection Rate : 0.59259
 Detection Prevalence : 0.96296
 Balanced Accuracy : 0.54545


'Positive' Class : Hamilton

Problem 4: 5-fold Cross Validation


- (a) For each value of $\mu \in \left\{ \frac{0.1}{|D|}, \frac{1}{|D|}, \frac{10}{|D|}, \frac{100}{|D|}, \frac{1000}{|D|} \right\}$, where $|D| = 4875$ is the size of our dictionary; estimations of the the correct classification rate, the false negative rate, and the false positive rate using 5-fold cross-validation are outlined below. For each of the metrics, the table indicates the value for each of the 25 tests and the graph indicates averages over the 5 tests for each choice of μ .

Correct Rate:

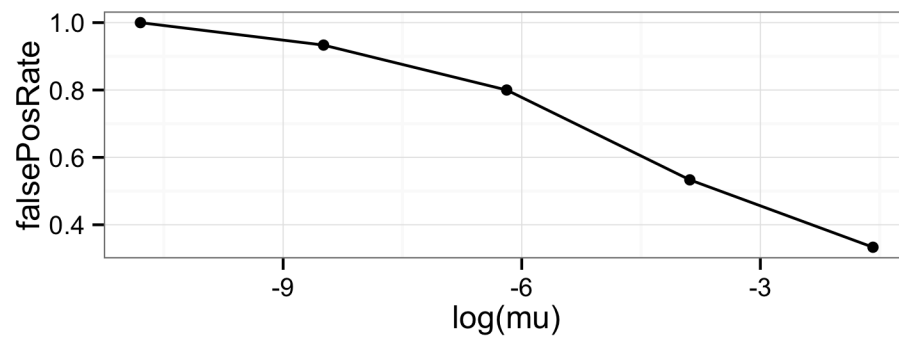
	fold1	fold2	fold3	fold4	fold5
mu1	0.7	0.7	0.7	0.7	0.7
mu2	0.8	0.7	0.7	0.7	0.7
mu3	0.8	0.7	0.7	0.8	0.8
mu4	0.9	0.7	0.8	0.8	1.0
mu5	1.0	0.8	0.9	0.8	1.0

**False Negative Rate:**

	fold1	fold2	fold3	fold4	fold5
mu1	0	0	0	0	0
mu2	0	0	0	0	0
mu3	0	0	0	0	0
mu4	0	0	0	0	0
mu5	0	0	0	0	0

**False Positive Rate:**

	fold1	fold2	fold3	fold4	fold5
mu1	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000
mu2	0.6666667	1.0000000	1.0000000	1.0000000	1.0000000
mu3	0.6666667	1.0000000	1.0000000	0.6666667	0.6666667
mu4	0.3333333	1.0000000	0.6666667	0.6666667	0.0000000
mu5	0.0000000	0.6666667	0.3333333	0.6666667	0.0000000



- (b) Based on the values measured on the training set (using 5-fold cross validation), the best value for μ is $\mu_5 = \frac{1000}{|D|} \approx 0.205$. At this value we maximize the accuracy and minimize the false positive rate, with no increase in the false negative rate.
- (c) For each value of μ , train on the full training set and test on the full testing set. Summarize the correct classification rate, the false negative rate and the false positive rate in three graphs. Does your answer from (b) still seem the best value? Why or why not?
- (d) How close are the rates estimated from cross-validation to the true rates on the testing set (give percentage error)? What could account for differences? Give one way the differences between the cross-validation rate estimates and the rates on the training sets could be minimized.